

A sentiment model trained from scratch

PyTorch classifier • UCI Sentiment Labelled Sentences • Evaluation on unseen test data

TEST ACCURACY

68.82%

309 correct predictions out of 449

PERCENTAGE POINTS

+18.71

Above the majority-class baseline

MACRO-F1

0.6811

F1 averaged across both classes

Accuracy on 449 test sentences

Both evaluated on the same held-out split

Custom PyTorch model



Majority-class baseline



0% 25% 50% 75% 100%

Baseline predicts negative for every test sentence.

Prediction breakdown

Confusion matrix • Number of sentences

	PREDICTED	
	Negative	Positive
ACTUAL		
Negative	188	37
Positive	103	121

Teal cells = correct predictions

103 positive sentences were missed.

HOW THE MODEL WAS BUILT

1,200 training sentences • Vocabulary fit on training only • Checkpoint chosen on validation

Token embeddings → masked mean pooling → hidden layer → two sentiment scores